

AI1072: Machine learning, exercise sheet 2

MNIST can be obtained via

<http://filr.hs-fulda.de/ssf/s/readFile/share/206/-8986560519763980414/publicL>

numpy& matplotlib with MNIST

Load the MNIST data as follows:

```
import gzip, pickle ;
with gzip.open("MNIST.pkl.gz") as f:
    data = pickle.load(f) ;
    traind = data["data_train"]
    trainl = data["labels_train"]
```

'traind' contains the samples, 'trainl' the target values (labels) in one-hot format.

- a) Inspect both arrays and especially print out their shape. How many samples and targets are there? How many numbers constitute a sample? Why are there 10 dimensions to a single target value?
- b) Slice out the 1000th sample and display it. Print its class as well!
- c) Display a histogram over the classes of all samples!
- d) Print the number of samples having class 5
- e) Display the classes of the first 100 samples as a bar plot!
- f) Generate the following variations of the 10th sample and display them together:
 - flip along x axis
 - flip along y axis
- g) Extract from the training data just the samples having class 4 and display the last 4 of them!
- h) Copy out 100 randomly chosen samples and display the first 3!

Matplotlib

- a) plot the sine function between -5 and 5 using 100 support points! Use `numpy.sin` or `math.sin` to generate the values!
- b) generate a scatter plot of the same data as in a)!
- c) generate a bar plot of the same data as in a)!
- d) plot sine and cosine together in a single plot